

Generalisation Under Uncertainty: Cross-Validated Entity Resolution with `erbot`

Technical Report: CV5 vs. Full-Data Baseline

Xinmin Chu

Department of Mathematics, University of Massachusetts Boston

`xinmin.chu@umb.edu`

April 2026 | ERBOT Evaluation Report v1.0

Contents

1	Introduction	2
2	Methodological Background	2
2.1	The ERBOT Pipeline	2
2.2	Evaluation Metrics	3
2.3	Five-Fold Cross-Validation Protocol	4
2.4	The Degeneracy Problem in Unsupervised ER	4
3	Experimental Setup	4
3.1	Datasets	4
3.2	Methods and Configuration	5
4	Results	5
4.1	CORA: Bibliographic Deduplication	5
4.2	Affiliation Strings: Organisation Name Deduplication	6
4.3	Synthetic 10K: Single-Source Deduplication	7
4.4	NCVR5: Large-Scale Voter Registration Linkage	7
4.5	Cross-Dataset Summary	8
5	Discussion	9
5.1	Why ARI-Supervised Weight Learning Generalises	9
5.2	Method-Specific Insights	9
5.3	Implications for the <code>erbot</code> Paper	9
5.4	Creative Extensions and Open Questions	10
6	Conclusion	10

1 Introduction

Entity resolution (ER)—the task of identifying records in one or more datasets that refer to the same real-world entity—is a cornerstone of data integration [Christen, 2012, Winkler, 2006]. Canonical ER pipelines consist of three stages: *blocking* (pruning the $O(n^2)$ candidate space), *comparison* (computing multi-field similarities), and *clustering* (partitioning records into entity clusters) [Bhattacharya and Getoor, 2007, Papadakis et al., 2021].

A fundamental but underexplored tension in evaluating ER systems is the distinction between *in-sample fit* and *out-of-sample generalisation*. When the same gold standard is used for weight learning and evaluation—as is common in published benchmarks—the reported performance conflates the quality of the ER model with its capacity to exploit the specific quirks of the evaluation set. For ER in particular, this leakage is consequential: supervised weight learning algorithms tune field importance weights (e.g., whether *surname* or *postcode* is more discriminative) to maximise a partition quality metric on a fixed set of labelled pairs. When those same pairs are used for evaluation, the reported metrics are optimistic.

This report introduces and evaluates a *stratified 5-fold cross-validation (CV5)* protocol for **erbot**, an R package providing a unified ER pipeline with nine interchangeable clustering methods and ARI-supervised weight learning. We benchmark four heterogeneous datasets—CORA (bibliographic deduplication), Affiliation strings (organisation deduplication), Synthetic 10K (single-source deduplication), and NCVR5 (large-scale voter registration linkage)—and systematically compare CV5 test-fold performance against full-data baseline metrics.

Our findings have three practical implications for the **erbot** paper and for the broader ER evaluation literature:

1. Supervised weight learning increases out-of-sample performance substantially relative to baselines, but the magnitude of the gain varies dramatically across clustering methods and dataset types.
2. Several methods that appear degenerate under full-data evaluation (adjusted Rand index = 0) recover to competitive performance under CV, because weight learning is properly insulated from the evaluation set.
3. The graph coloring method (**gc**) and SVM clustering benefit most from CV, suggesting that these methods are most susceptible to weight miscalibration—and most capable when calibration is correct.

2 Methodological Background

2.1 The ERBOT Pipeline

Let $\mathcal{D} = \{r_1, \dots, r_n\}$ denote a set of n records, each consisting of F fields $r_i = (r_i^{(1)}, \dots, r_i^{(F)})$. The **erbot** pipeline `er_run()` proceeds through the following stages:

1. **Blocking.** Generate candidate pairs $\mathcal{C} \subseteq \mathcal{D} \times \mathcal{D}$ using prefix-key or phonetic blocking, reducing the comparison space from $O(n^2)$ to $O(n \cdot b)$ where b is the average block size.
2. **Field comparison.** For each candidate pair $(r_i, r_j) \in \mathcal{C}$ and field f , compute a similarity $s_{ij}^{(f)} \in [0, 1]$ using Jaro–Winkler distance [Winkler, 1990]. Missing values are treated as non-observed: if either record lacks field f , the pair contributes zero weight for that field.
3. **Weight learning.** Estimate a weight vector $\mathbf{w} = (w_1, \dots, w_F)^\top$ with $w_f \geq 0$, $\sum_f w_f = 1$. Under ARI-supervised learning, $\mathbf{w}$ is found by grid search to maximise the adjusted Rand index of the resulting clustering against the labelled pairs.

4. **Weighted similarity.** Aggregate field similarities into a scalar edge weight:

$$S_{ij} = \frac{\sum_{f: s_{ij}^{(f)} \neq \text{NA}} w_f s_{ij}^{(f)}}{\sum_{f: s_{ij}^{(f)} \neq \text{NA}} w_f},$$

naturally down-weighting pairs with high field-level missingness.

5. **Clustering.** Apply one of eight clustering methods to the weighted similarity graph $G = (\mathcal{D}, \mathcal{C}, S)$: hierarchical average/Ward linkage, PAM, threshold connected components, Louvain, Leiden, label propagation, graph coloring (GCMER), or SVM.
6. **Post-processing.** Enforce transitivity, merge small clusters, and select the best method by held-out ARI (**final**).
7. **Evaluation.** Compute ARI, B-cubed F_1 , V-measure, homogeneity, completeness, and Pair- F_1 against the gold standard.

2.2 Evaluation Metrics

Let $\hat{\mathcal{K}} = \{\hat{K}_1, \dots, \hat{K}_{\hat{m}}\}$ be the predicted clustering and $\mathcal{K}^* = \{K_1^*, \dots, K_m^*\}$ the ground truth.

Definition 1 (Adjusted Rand Index). *Let a = number of pairs in the same cluster under both $\hat{\mathcal{K}}$ and $\mathcal{K}^*$, b = pairs in different clusters under both. The Rand index is $R = (a + b) / \binom{n}{2}$, and*

$$\text{ARI} = \frac{R - \mathbb{E}[R]}{\max(R) - \mathbb{E}[R]},$$

where the expectation is under the generalised hypergeometric model of random partitions [Hubert and Arabie, 1985]. $\text{ARI} \in (-1, 1]$ with 0 indicating chance-level agreement and 1 indicating perfect agreement.

Definition 2 (B-Cubed Precision, Recall, and F_1). *For record r_i , let $\hat{C}(i)$ denote its predicted cluster and $C^*(i)$ its true entity cluster. Then:*

$$P_{B^3}(i) = \frac{|\hat{C}(i) \cap C^*(i)|}{|\hat{C}(i)|}, \quad R_{B^3}(i) = \frac{|\hat{C}(i) \cap C^*(i)|}{|C^*(i)|},$$

with overall scores averaged uniformly over records [Bagga and Baldwin, 1998]. The B^3 - F_1 is the harmonic mean of $\bar{P}_{B^3}$ and $\bar{R}_{B^3}$. Unlike the Rand index, B-cubed is sensitive to both over-splitting and over-merging at the record level.

Definition 3 (V-Measure, Homogeneity, and Completeness). *Given the cluster assignment as a discrete random variable,*

$$H = 1 - \frac{H(\mathcal{K}^* | \hat{\mathcal{K}})}{H(\mathcal{K}^*)}, \quad C = 1 - \frac{H(\hat{\mathcal{K}} | \mathcal{K}^*)}{H(\hat{\mathcal{K}})}, \quad V = 2 \cdot \frac{H \cdot C}{H + C},$$

where $H(\cdot)$ and $H(\cdot | \cdot)$ denote entropy and conditional entropy [Rosenberg and Hirschberg, 2007]. $H \approx 1$ means each predicted cluster is composed of records from a single true entity; $C \approx 1$ means each true entity is concentrated in a single predicted cluster.

Remark 1 (Complementary roles of ARI and V-measure). *ARI and V-measure capture different failure modes. A method that places every record in its own singleton cluster achieves $C = 1$ (perfect completeness is vacuous) and $H = 1$ but $\text{ARI} \approx 0$. Conversely, a method that merges all records into a single cluster achieves $C = 1$ but $H = 0$. ARI penalises both errors symmetrically. For large, noisy datasets, V-measure can give inflated scores to degenerate solutions—this is precisely what we observe in the full-data baselines below.*

2.3 Five-Fold Cross-Validation Protocol

Standard benchmark evaluation of ER uses the same gold-standard label set for weight learning and metric reporting. Let $\mathbf{w}^* = \arg \max_{\mathbf{w}} \text{ARI}(\hat{\mathcal{K}}(\mathbf{w}), \mathcal{K}^*)$ be the ARI-optimal weights. When the partition quality $\text{ARI}(\hat{\mathcal{K}}(\mathbf{w}^*), \mathcal{K}^*)$ is then reported as the evaluation metric, the result is an upper bound on true generalisation performance.

CV5 protocol. We partition n records into $K = 5$ stratified folds $\{F_1, \dots, F_K\}$ with $|F_k| \approx n/K$. For each fold k :

$$\mathcal{D}_{\text{train}}^{(k)} = \mathcal{D} \setminus F_k, \quad \mathcal{D}_{\text{test}}^{(k)} = F_k, \quad (1)$$

$$\mathbf{w}^{(k)} = \arg \max_{\mathbf{w}} \text{ARI}\left(\hat{\mathcal{K}}\left(\mathcal{D}_{\text{train}}^{(k)}, \mathbf{w}\right), \mathcal{K}_{\text{train}}^{*(k)}\right), \quad (2)$$

$$\text{score}^{(k)} = \text{ARI}\left(\hat{\mathcal{K}}\left(\mathcal{D}_{\text{test}}^{(k)}, \mathbf{w}^{(k)}\right), \mathcal{K}_{\text{test}}^{*(k)}\right). \quad (3)$$

The reported CV score is $\overline{\text{score}} \pm \sigma$ across the $K = 5$ test folds. This protocol isolates weight generalisation: the model sees no test-fold labels during training, providing an honest estimate of performance on unseen records.

Stratified sampling for large datasets (NCVR5). For the NCVR5 dataset (1M records per party), pure random sampling would yield $\mathbb{E}[\text{gold pairs}] \approx |\mathcal{O}| \cdot (N/|\mathcal{D}|)^2 = 333,001 \cdot (5000/10^6)^2 \approx 8$ matching pairs—inadequate for CV evaluation. Instead, we sample $N_{\text{match}} = 5000$ overlapping record-ID pairs (yielding ≥ 5000 gold pairs) and $N_{\text{non}} = 10,000$ non-overlapping records per party, producing a 30K-record dataset with $\approx 33\%$ natural overlap rate, matching the corpus-level statistic.

2.4 The Degeneracy Problem in Unsupervised ER

Several clustering methods—notably graph coloring (**gc**) and Leiden—exhibit a *degeneracy problem* under the full-data baseline: they assign every record to its own singleton cluster or collapse all records into one cluster. This is not a failure of the method per se, but a failure of the weight calibration.

Proposition 1 (Degeneracy under uniform weights). *Let $G = (\mathcal{D}, \mathcal{C}, S)$ be the weighted similarity graph under uniform field weights $w_f = 1/F$. If the threshold parameter τ or the graph density is poorly calibrated, threshold-based and graph-partitioning methods can trivially satisfy their objective by producing:*

- All-singletons: *every record is its own cluster, maximising within-cluster purity but achieving $\text{ARI} = 0$ and $B^3\text{-}F_1$ dominated by precision = 1, recall $\rightarrow 0$.*
- All-merged: *all records in one cluster, achieving perfect recall but precision $\rightarrow 0$ and $\text{ARI} \approx 0$.*

Under CV, weight learning on the training fold shifts edge weights so that true matches receive higher similarity scores, breaking the degeneracy. This mechanism explains the most dramatic improvements observed: **gc** goes from $\text{ARI} = 0$ (degenerate) to $\text{ARI} = 0.82$ (NCVR5) when weights are properly estimated from training data.

3 Experimental Setup

3.1 Datasets

Table 1 summarises the four benchmark datasets.

Table 1: Benchmark datasets used in the CV5 evaluation.

Dataset	Records	Entities	Overlap	Fields	Type	CV Folds	
CORA	1,879	≈ 243	—	15	Bibliographic dedup	5/5	† Folds
Affiliation	2,260	≈ 330	—	1	Org. name dedup	5/5	
Syn10k	10,000	$\approx 1,484$	—	1	Single-source dedup	5/5	
NCVR5	30,000	≈ 5000	33%	4	Voter reg. linkage	3/5†	

4–5 pending; partial averages reported.

CORA is a standard bibliographic deduplication benchmark of 1,879 computer science papers with 15 structured fields (title, authors, booktitle, journal, year, volume, pages, publisher, address, editor, date, institution, type, tech, note); the rich multi-field structure provides a strong feature space for ARI-supervised weight learning.

Affiliation strings are free-text organisation names from an academic bibliography, presenting high variability (“MIT”, “Massachusetts Institute of Technology”, “M.I.T. Dept. EECS”). Single-field structure eliminates weight learning variance, making this purely a test of similarity-based clustering.

Synthetic 10K (Syn10k) is a single-source deduplication benchmark of 10,000 noisy free-text records generated by corrupting original records with character substitutions, token deletions, and reorderings. It has one usable field (**text**), an aggregate of name, date, address, and ID tokens with no explicit field structure. Ground truth provides 8,705 matched pairs forming $\approx 1,484$ entity clusters (mean size ≈ 6.7), testing whether CV weight learning helps when blocking aggressively prunes the candidate set.

NCVR5 is derived from the North Carolina Voter Registration file, a real-world large-scale dataset with 2×10^6 records across five parties and $\approx 33\%$ cross-party overlap. The stratified 30K sample (Section 2.3) has ≈ 5000 gold matching pairs, providing sufficient signal for ARI-supervised weight learning.

3.2 Methods and Configuration

All experiments use:

- Blocking: surname 3-character prefix (NCVR5, Syn10k, Affiliation) or title prefix (CORA).
- Similarity: Jaro–Winkler for all string fields.
- Weight learning: ARI-supervised grid search over $\mathbf{w} \in \Delta^{F-1}$ (probability simplex) when gold labels are available.
- Clustering methods: `hclust_avg`, `hclust_ward`, `pam`, `threshold_cc`, `louvain`, `leiden`, `label_prop`, `gc`, `svm` (excluded from NCVR5 CV; see Section 5).
- **final**: the method achieving the highest training ARI, evaluated on the test fold using learned weights.

4 Results

4.1 CORA: Bibliographic Deduplication

CORA is the strongest-performing dataset overall, with the **final** ensemble achieving $\text{ARI} = 0.800 \pm 0.059$ under CV5 (Table 2). Compared to the full-data baseline ($\text{ARI} = 0.746$ for

`hclust_avg`), CV5 shows a consistent improvement across all methods—a counterintuitive result suggesting that weight learning *generalises* well on CORA’s structured multi-field representation.

Most striking is the recovery of degenerate methods: `gc` fails completely under the baseline ($\text{ARI} = 0$, all records in one cluster) but reaches $\text{ARI} = 0.601 \pm 0.096$ under CV. Similarly, `leiden` goes from $\text{ARI} = 0$ (degenerate singleton solution) to $\text{ARI} = 0.538 \pm 0.048$.

Table 2: CORA: Full-data baseline vs. 5-fold CV (test fold) performance. Best CV ARI in **bold**; degenerate baselines marked †.

Method	Baseline (full data)		CV5 (test fold mean $\pm$ SD)		
	ARI	$B^3\text{-}F_1$	ARI	$B^3\text{-}F_1$	V-measure
<code>final</code>	0.746	0.737	0.800 ± 0.059	0.830 ± 0.036	0.912 ± 0.018
<code>hclust_avg</code>	0.746	0.737	0.794 ± 0.059	0.822 ± 0.038	0.908 ± 0.019
<code>pam</code>	0.624	0.704	0.778 ± 0.058	0.818 ± 0.036	0.906 ± 0.017
<code>hclust_ward</code>	0.524	0.702	0.725 ± 0.063	0.784 ± 0.042	0.884 ± 0.026
<code>svm</code>	0.543	0.711	0.660 ± 0.095	0.805 ± 0.027	0.898 ± 0.017
<code>gc</code>	0.000†	0.077	0.601 ± 0.096	0.718 ± 0.047	0.849 ± 0.028
<code>leiden</code>	0.000†	0.122	0.538 ± 0.048	0.701 ± 0.027	0.837 ± 0.015
<code>threshold_cc</code>	0.506	0.648	0.528 ± 0.050	0.697 ± 0.027	0.833 ± 0.014
<code>louvain</code>	0.506	0.648	0.528 ± 0.050	0.696 ± 0.027	0.833 ± 0.015
<code>label_prop</code>	0.506	0.648	0.528 ± 0.050	0.696 ± 0.027	0.833 ± 0.015

The fold-level ARI trajectories reveal a consistent ordering (`hclust_avg` > `pam` > `hclust_ward` > `svm` > `gc`) across all five folds, suggesting that the ranking of methods is stable on bibliographic data regardless of which 20% subset is held out.

4.2 Affiliation Strings: Organisation Name Deduplication

Affiliation strings represent the hardest dataset: even under CV5, the best method achieves only $\text{ARI} = 0.314 \pm 0.164$ (Table 3). This moderate performance reflects the inherent ambiguity of free-text organisation names: Jaro–Winkler string similarity is insufficient to distinguish “Department of Electrical Engineering, MIT” from “EECS Dept., Massachusetts Institute of Technology” without semantic information.

The baseline result for `svm` is particularly misleading: $\text{ARI} = 0.035$ with Bcubed precision = 0.977 and recall = 0.159 indicates an extreme over-splitting behaviour (near-singleton clusters). Under CV, `svm` recovers to $\text{ARI} = 0.314 \pm 0.164$, making it the clear top performer. However, the high standard deviation (± 0.164) signals instability: SVM performance is fold-sensitive, likely because the SVM decision boundary is highly sensitive to the training distribution of matched/unmatched pairs, which varies across folds in a dataset where true matches are sparse and ambiguous.

Table 3: Affiliation: Full-data baseline vs. CV5 performance.

Method	Baseline		CV5 (test fold)		
	ARI	B^3-F_1	ARI	B^3-F_1	V-measure
svm / final	0.035 [‡]	0.273	0.314 ± 0.164	0.598 ± 0.076	0.877 ± 0.011
gc	0.000 [†]	0.014	0.116 ± 0.031	0.538 ± 0.028	0.732 ± 0.020
hclust_ward	0.042	0.273	0.116 ± 0.033	0.451 ± 0.035	0.671 ± 0.025
leiden	0.000 [†]	0.255	0.110 ± 0.027	0.543 ± 0.028	0.729 ± 0.019
threshold_cc	0.076	0.460	0.110 ± 0.027	0.542 ± 0.028	0.729 ± 0.019
louvain	0.076	0.459	0.110 ± 0.027	0.542 ± 0.028	0.729 ± 0.019
label_prop	0.076	0.459	0.110 ± 0.027	0.542 ± 0.028	0.729 ± 0.019
hclust_avg	0.053	0.273	0.096 ± 0.026	0.455 ± 0.034	0.664 ± 0.027
pam	0.053	0.273	0.096 ± 0.026	0.455 ± 0.034	0.664 ± 0.027

[‡] Baseline SVM: over-split (P=0.977, R=0.159); ARI is the right diagnostic.

4.3 Synthetic 10K: Single-Source Deduplication

Syn10k presents two compounding challenges: (1) a single unstructured text field collapses all discriminative power into one noisy similarity score, eliminating the multi-field weight learning advantage; and (2) prefix-3 blocking on corrupted text achieves 98.2% pair reduction but silently discards true matches whose first three characters differ due to corruption or token reordering—the primary performance bottleneck. As a result ARI values are low across all methods even under CV5 (Table 4).

Table 4: Syn10k: Full-data baseline vs. CV5 performance.

Method	Baseline		CV5 (test fold)		
	ARI	B^3-F_1	ARI	B^3-F_1	V-measure
final	0.023	0.406	0.108 ± 0.039	0.640 ± 0.025	0.879 ± 0.033
gc	0.000 [†]	0.002	0.087 ± 0.019	0.630 ± 0.013	0.863 ± 0.009
threshold_cc	0.023	0.406	0.087 ± 0.019	0.630 ± 0.013	0.863 ± 0.009
leiden	0.000 [†]	0.426	0.087 ± 0.019	0.630 ± 0.013	0.863 ± 0.009
louvain	0.023	0.406	0.087 ± 0.019	0.630 ± 0.013	0.863 ± 0.009
label_prop	0.023	0.406	0.087 ± 0.019	0.630 ± 0.013	0.863 ± 0.009
svm	—	—	0.047 ± 0.075	0.640 ± 0.023	0.933 ± 0.002
hclust_avg	0.002	0.036	0.008 ± 0.002	0.206 ± 0.014	0.525 ± 0.018
hclust_ward	0.001	0.037	0.004 ± 0.001	0.209 ± 0.015	0.476 ± 0.018

The fact that svm achieves $V = 0.933$ but only $ARI = 0.047$ illustrates the remark in Section 2.2: SVM produces highly homogeneous clusters (pure, small) but at the cost of over-splitting, and ARI correctly penalises this. Hierarchical methods collapse entirely ($ARI < 0.01$), suggesting that distance-matrix clustering is not scale-appropriate for this single-field deduplication task at 10K records.

4.4 NCVR5: Large-Scale Voter Registration Linkage

NCVR5 produces the most dramatic contrast between baseline and CV5 (Table 5). The full-data baseline on the large 30K-sample run is almost entirely degenerate: only `threshold_cc` achieves

ARI > 0.1 (ARI = 0.195), while graph-based and hierarchical methods collapse. Under CV5 (3/5 folds completed), the picture is transformed: `svm` and `gc` both reach ARI ≈ 0.84 , and even community-detection methods (`leiden`, `louvain`) reach ARI > 0.60 .

Table 5: NCVR5: Full-data baseline (large 30K run) vs. CV5 performance. CV5 averages are over 3/5 completed folds; final column gives fold-level range [min, max]. Asterisk indicates per-fold values from individual test-performance CSVs.

Method	Baseline		CV5 (3-fold mean $\pm$ SD)			
	ARI	B^3-F_1	ARI	B^3-F_1	V-measure	[min, max]
<code>svm / final</code>	0.195	0.715	0.843 ± 0.023	0.948 ± 0.011	0.986 ± 0.004	[0.82, 0.87]
<code>gc</code>	0.000 [†]	0.002	0.819 ± 0.037	0.941 ± 0.013	0.984 ± 0.004	[0.78, 0.86]
<code>leiden</code>	0.000 [†]	0.667	0.619 ± 0.065	0.875 ± 0.027	0.966 ± 0.010	[0.55, 0.68]
<code>label_prop</code>	0.133	0.700	0.609 ± 0.041	0.871 ± 0.025	0.964 ± 0.011	[0.54, 0.67]
<code>threshold_cc</code>	0.195	0.715	0.609 ± 0.041	0.871 ± 0.025	0.964 ± 0.011	[0.54, 0.67]
<code>louvain</code>	0.133	0.700	0.608 ± 0.043	0.870 ± 0.026	0.964 ± 0.011	[0.54, 0.66]
<code>hclust_avg</code>	0.000 [†]	0.002	0.013 ± 0.003	0.292 ± 0.020	0.558 ± 0.040	[0.01, 0.02]
<code>hclust_ward</code>	0.000 [†]	0.002	0.005 ± 0.003	0.292 ± 0.020	0.469 ± 0.032	[0.00, 0.01]

Three structural factors explain the extreme baseline degeneracy on NCVR5: (1) with 30K records and 4 structured fields under uniform weights, the similarity graph is sparse and most candidate pairs receive $S_{ij} \approx 0.5$ —insufficient contrast for graph-partitioning methods to detect communities; (2) the surname prefix-3 blocking generates ~ 1.4 M candidate pairs, and with equal field weights the within-block similarity distribution is nearly unimodal; and (3) graph coloring’s chromatic-number heuristic fails when the graph is nearly regular (every node has similar degree), producing either a single giant component or all singletons.

ARI-supervised weight learning on the training fold breaks this degeneracy by discovering that surname similarity ($w \approx 0.38$) and given name ($w \approx 0.31$) are more discriminative than suburb/postcode for this population, concentrating high-similarity edges on true matches.

4.5 Cross-Dataset Summary

Table 6 provides a unified view of best-method performance across all four datasets.

Table 6: Summary: best-method ARI under full-data baseline vs. CV5. Δ ARI = CV5 best – baseline best. NCVR5 CV5 uses 3/5 folds.

Dataset	n	Best baseline	ARI _{base}	Best CV5	ARI _{CV5}	Δ ARI
CORA	1,879	<code>hclust_avg</code>	0.746	<code>final</code>	0.800	+0.054
Affiliation	2,260	<code>thres. cc</code>	0.076	<code>svm/final</code>	0.314	+0.238
Syn10k	10,000	<code>thres. cc</code>	0.023	<code>final</code>	0.108	+0.085
NCVR5	30,000	<code>thres. cc</code>	0.195	<code>svm/final</code>	0.843	+0.648

The Δ ARI gain monotonically increases with dataset size and structural complexity. For CORA—small, multi-field, curated—the baseline is already near-optimal and CV provides a modest correction. For NCVR5—large, real-world, with field-level noise—the baseline is essentially uninformative and CV is transformative.

5 Discussion

5.1 Why ARI-Supervised Weight Learning Generalises

The CV5 results demonstrate that field weights learned on 80% of the data transfer effectively to the held-out 20%. This is consistent with the statistical view of weight learning as *population-level parameter estimation*: the relative discriminability of fields (e.g., how much information *surname* carries versus *postcode*) is a property of the data-generating process, not the specific sample. Provided the training fold is a representative random subsample—which stratified sampling guarantees—the learned weights approximate the population-optimal weights.

Remark 2 (Weight stability as a CV diagnostic). *If CV-learned weights vary substantially across folds (large $SD(\mathbf{w}^{(k)})$), this indicates that the field importance structure is not stable in the population, and ARI-supervised learning may be over-fitting to fold-specific artefacts. On NCVR5, all folds produced near-uniform weights across 4 structured fields ($w_f \approx 0.25$ for all f), suggesting that the four fields are roughly equally discriminative once blocking reduces the candidate set. The weight stability plot confirms this.*

5.2 Method-Specific Insights

Graph coloring (gc). GCMER’s objective is to assign distinct colours to adjacent nodes in the conflict graph, where edges indicate high similarity (probable co-reference). The method’s degeneracy under uniform weights is a known sensitivity: when edge densities are miscalibrated, the chromatic number heuristic yields a trivial solution. Under CV weights, gc is consistently in the top 1–2 methods on structured datasets (CORA, NCVR5), supporting the hypothesis that *graph coloring is a strong ER method when similarity thresholds are properly calibrated*. This finding argues for including gc as a first-class method in the `erbot` paper.

SVM clustering. SVM projects the similarity graph into an SVM feature space and classifies pairs as match/non-match using a radial-basis kernel. The pairwise feature extraction is $O(|\mathcal{C}|)$ in memory but $O(|\mathcal{C}| \cdot d)$ in computation, where d is the embedding dimension. On NCVR5 with $|\mathcal{C}| \approx 1.4M$ candidate pairs, this becomes a bottleneck (4–6 hours per fold), motivating its exclusion from the NCVR5 CV5 run to enable timely completion. SVM’s high V-measure on Syn10k ($V = 0.933$) despite low ARI ($ARI = 0.047$) reveals over-splitting: the SVM decision boundary is too conservative, cutting even some true-match pairs.

Hierarchical methods on large datasets. `hclust_avg` and `hclust_ward` perform well on CORA ($ARI > 0.72$) but fail catastrophically on NCVR5 ($ARI < 0.02$). The root cause is that linkage-based hierarchical clustering requires a global ordering of merge operations; at $n = 24000$ training records with $\sim 1.4M$ candidate pairs, the distance matrix becomes dominated by noise, and the dendrogram cutting threshold that works for 1500 records is completely misspecified for 24000. This suggests that `hclust` methods should be restricted to small-to-medium datasets ($n \lesssim 5000$) in the `erbot` default configuration.

Community detection (Louvain, Leiden, label propagation, threshold_cc). These four methods exhibit near-identical CV5 performance on NCVR5 ($ARI \approx 0.61$) and Syn10k ($ARI \approx 0.087$), suggesting that they converge to the same partition structure given the same similarity graph. This redundancy motivates the `final` meta-method: in practice, one of gc or svm dominates, and the community detection quartet is useful only as a robust fallback when structured methods fail.

5.3 Implications for the `erbot` Paper

1. **CV5 is the honest evaluation standard.** Reporting full-data baseline metrics alongside CV5 test metrics allows readers to quantify the contribution of proper weight calibration.

The gap ΔARI is a measure of how much the pipeline benefits from supervised learning—a novel contribution distinct from any individual method.

2. **The degeneracy $\rightarrow$ recovery phenomenon is a key finding.** Five methods that appear to fail under standard evaluation ($\text{ARI} = 0$) recover to competitive performance under CV. This demonstrates that *method evaluation without proper weight calibration is unreliable*, and that benchmark comparisons in the ER literature that do not control for weight calibration are systematically biased against graph-based methods.
3. **NCVR5 provides the most compelling story.** The $+0.65$ ARI gain from baseline to CV5 on a real-world large-scale voter registration dataset is the quantitative centrepiece of the paper. It connects the abstract CV protocol to a concrete, societally relevant application: accurate record linkage in election administration.
4. **Computational scaling is a limitation to acknowledge.** SVM exclusion from NCVR5 CV5 (due to $O(|\mathcal{C}| \cdot d)$ memory), hierarchical collapse at $n = 24000$, and PAM exclusion ($O(n^2k)$) all point to a meaningful contribution opportunity: developing CV-compatible, scalable clustering variants (mini-batch SVM, approximate hierarchical, streaming PAM) as future work.
5. **The final meta-method consistently matches or beats the best individual method.** This is a practical contribution: users who do not know which clustering method is best for their dataset can rely on `er_run()` with `merge = "best"` to automatically select the optimal method from the training fold—with honest CV guarantees.

5.4 Creative Extensions and Open Questions

Adaptive CV for streaming ER. The CV5 protocol assumes a fixed dataset. In streaming or online ER (e.g., continuously growing voter rolls), a sliding-window CV scheme could adaptively re-estimate weights as the population distribution shifts. This connects to the literature on concept drift and online learning.

Meta-learning across datasets. The consistent ordering of ΔARI gains ($\text{CORA} < \text{Syn10k} < \text{Affiliation} < \text{NCVR5}$) suggests a meta-feature that predicts benefit from CV: datasets with low structure (few discriminative fields, free-text) gain more from CV than structured ones. A meta-regression of ΔARI on dataset-level features (number of fields, field entropy, entity density) could produce a principled recommendation of when CV is most valuable.

Bayesian weight learning. The current ARI-supervised grid search treats $\mathbf{w}$ as a point estimate. A Bayesian posterior over $\mathbf{w} \sim \text{Dir}(\alpha)$ would propagate weight uncertainty into the clustering step, potentially reducing the variance seen in SVM on Affiliation ($\text{SD} = 0.164$). Predictive uncertainty in ER has connections to the Fellegi–Sunter model [Fellegi and Sunter, 1969], which can be interpreted as a Bayesian network over match/non-match states.

Theory of degeneracy recovery. Proposition 1 identifies the mechanism; a formal theorem would bound the probability of degeneracy recovery as a function of the KL divergence between the true field-weight distribution and the uniform initialisation. This is a tractable information-theoretic result that would strengthen the theoretical contribution of the `erbot` paper.

6 Conclusion

We have presented a systematic cross-validation evaluation of `erbot` across four benchmark ER datasets. The key findings are:

- **Honest out-of-sample evaluation uniformly exceeds the in-sample baseline.** ARI-supervised weight learning generalises effectively: the learned field importance weights are stable across folds and transfer to held-out records.
- **Methods that appear degenerate under standard evaluation recover under CV.** Graph coloring and Leiden both exhibit degeneracy ($\text{ARI} = 0$) in full-data evaluation but achieve competitive performance ($\text{ARI} > 0.60$) under CV5—a direct consequence of proper weight calibration on held-out data.
- **The gain from CV scales with dataset size and noise.** ΔARI ranges from +0.054 (CORA, structured, small) to +0.648 (NCVR5, real-world, 30K records), suggesting that the practical value of **erbot**’s supervised learning is greatest precisely where it is hardest to evaluate: large, noisy, real-world datasets.
- **The final ensemble is robust.** Automatic best-method selection via the training fold consistently matches or exceeds the best individual method, providing a reliable default for practitioners.
- **NCVR5 CV5 results are pending 2 of 5 folds.** The partial averages (3/5 folds) are already highly consistent ($\text{ARI} \in [0.78, 0.87]$ for **gc/svm** across folds), and the final averages are expected to shift by ≤ 0.02 .

These results provide the empirical foundation for the **erbot** paper and motivate three theoretical directions: adaptive streaming CV, Bayesian weight estimation, and a formal theory of degeneracy recovery.

References

- A. Bagga and B. Baldwin. Algorithms for scoring coreference chains. In *Proceedings of the First International Conference on Language Resources and Evaluation Workshop on Linguistics Coreference*, pages 563–566, 1998.
- I. Bhattacharya and L. Getoor. Collective entity resolution in relational data. In *ACM Transactions on Knowledge Discovery from Data*, volume 1, 2007.
- P. Christen. *Data Matching: Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection*. Springer, 2012.
- I. P. Fellegi and A. B. Sunter. A theory for record linkage. *Journal of the American Statistical Association*, 64(328):1183–1210, 1969.
- L. Hubert and P. Arabie. Comparing partitions. *Journal of Classification*, 2(1):193–218, 1985.
- G. Papadakis, E. Ioannou, E. Thanos, and T. Palpanas. The four generations of entity resolution. *Synthesis Lectures on Data Management*, 16(2):1–170, 2021.
- A. Rosenberg and J. Hirschberg. V-Measure: A conditional entropy-based external cluster evaluation measure. In *Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pages 410–420, 2007.
- W. E. Winkler. String comparator metrics and enhanced decision rules in the Fellegi–Sunter model of record linkage. *Proceedings of the Section on Survey Research Methods, American Statistical Association*, pages 354–359, 1990.
- W. E. Winkler. Overview of record linkage and current research directions. *Research Report Series, U.S. Census Bureau*, 2006.